

Processes and Threads

Chapter 3 and 4

*Operating Systems:
Internals and Design Principles, 9/E*
William Stallings

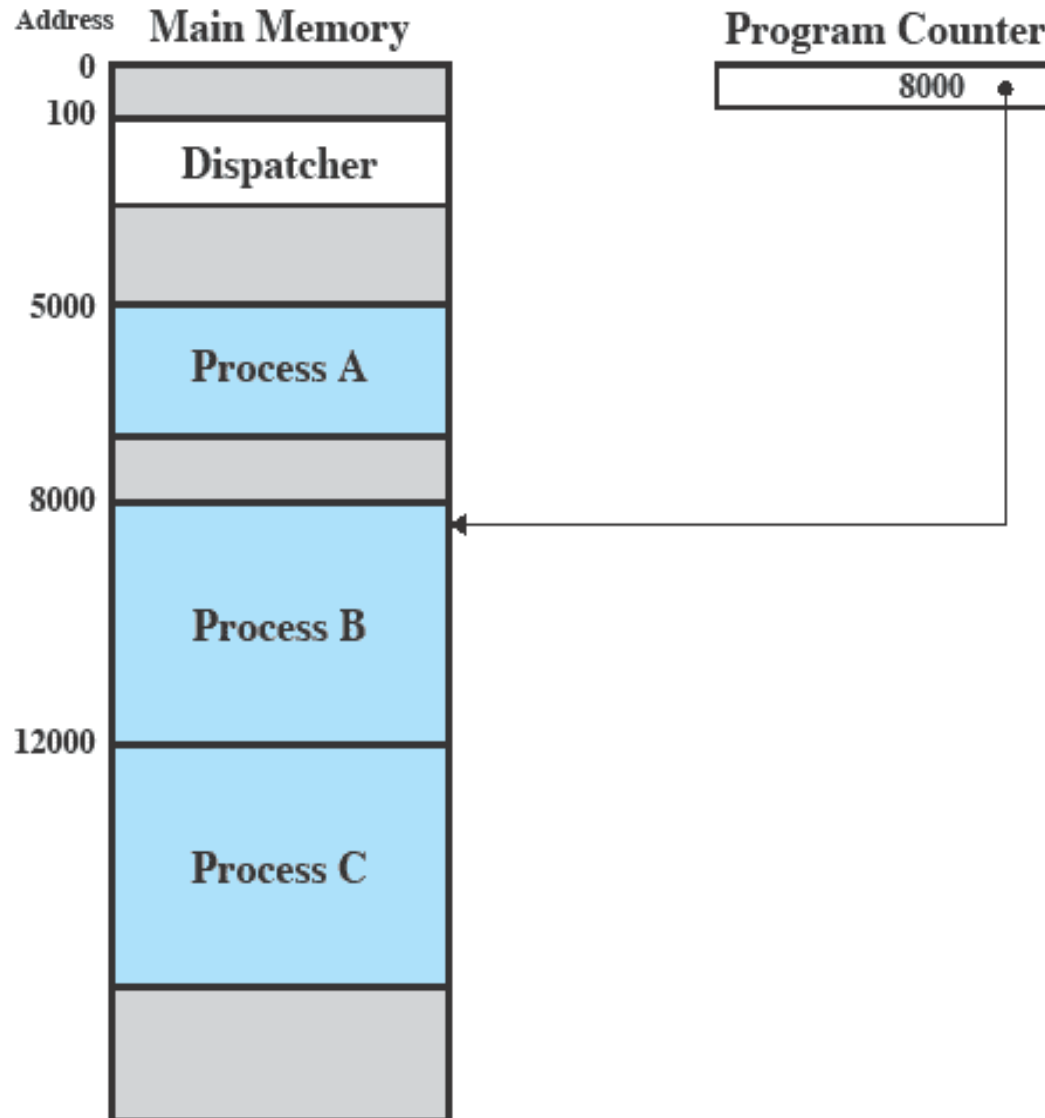
Process

- A program in execution
- An instance of a program running on a computer (cooking vs. recipe)
- The entity that can be assigned to and executed on a processor
- A unit of activity characterized by the **execution** of a sequence of instructions, a current **state**, and an associated set of system **resources**

Requirements of an Operating System

- Interleave the execution of multiple processes – **multiprogramming** -, to maximize processor utilization while providing reasonable response time
- Allocate resources to processes
- Support interprocess communication and user creation of processes

Example Execution



Combined Trace of Processes

1	5000		27	12004	
2	5001		28	12005	
3	5002				----- Timeout
4	5003		29	100	
5	5004		30	101	
6	5005		31	102	
		----- Timeout	32	103	
7	100		33	104	
8	101		34	105	
9	102		35	5006	
10	103		36	5007	
11	104		37	5008	
12	105		38	5009	
13	8000		39	5010	
14	8001		40	5011	
15	8002				----- Timeout
16	8003		41	100	
		----- I/O Request	42	101	
17	100		43	102	
18	101		44	103	
19	102		45	104	
20	103		46	105	
21	104		47	12006	
22	105		48	12007	
23	12000		49	12008	
24	12001		50	12009	
25	12002		51	12010	
26	12003		52	12011	
					----- Timeout

Interleaving


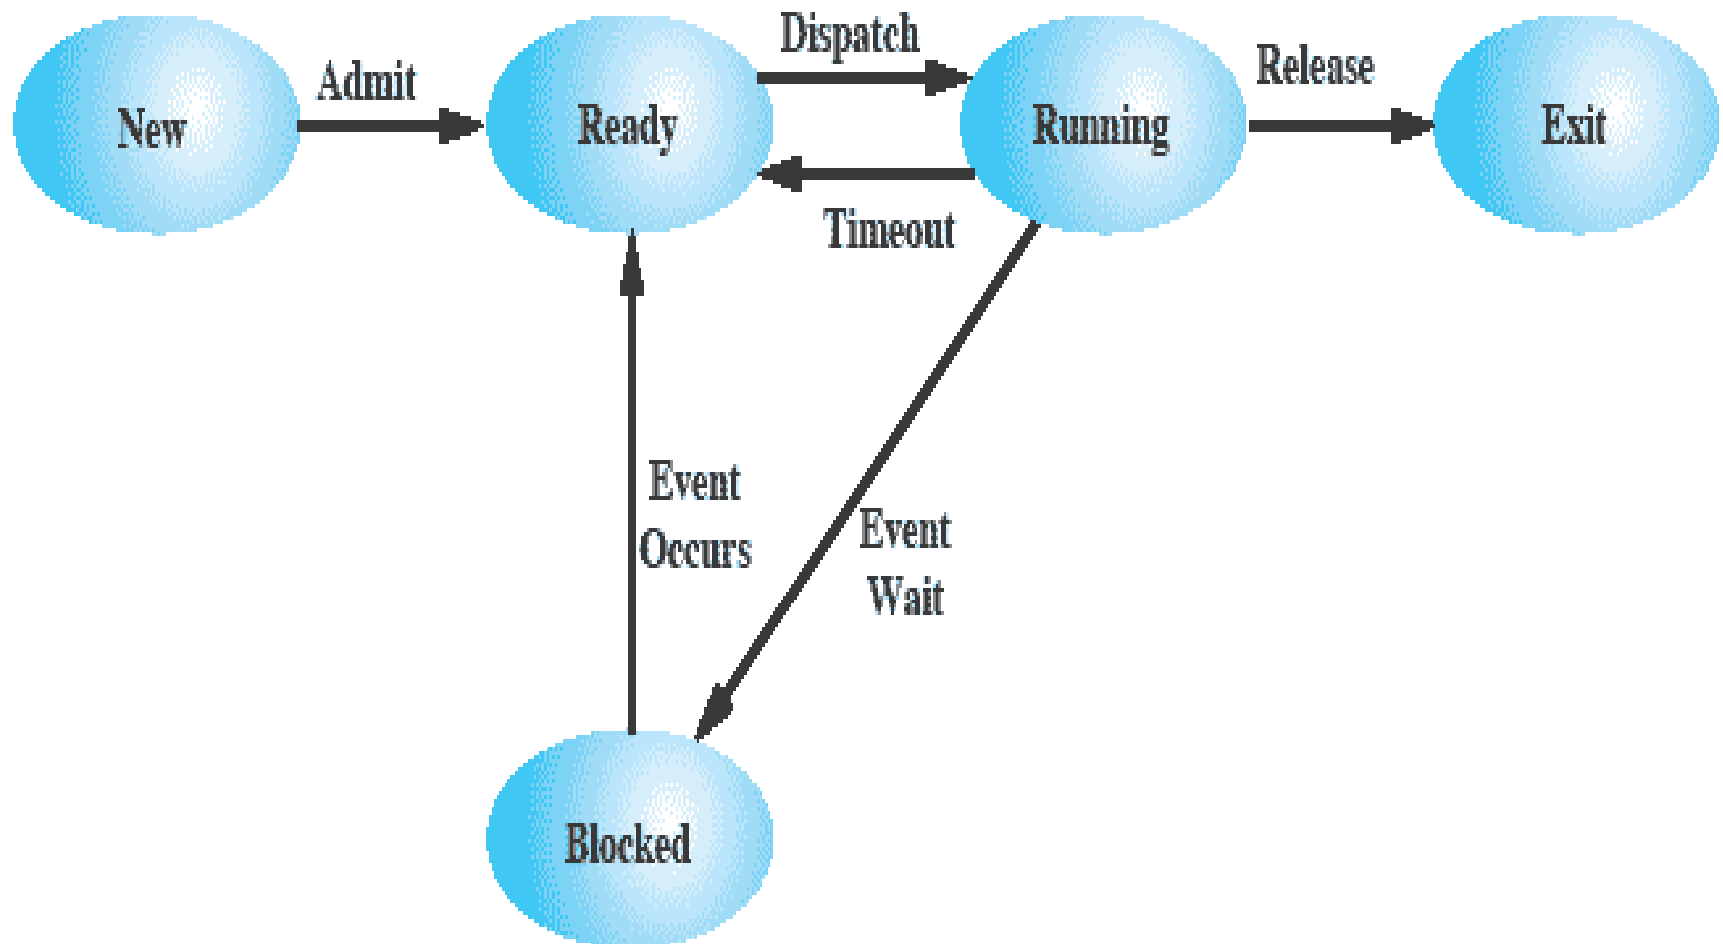


Figure 3.7 Process States for Trace of Figure 3.4

Five-State Process Model



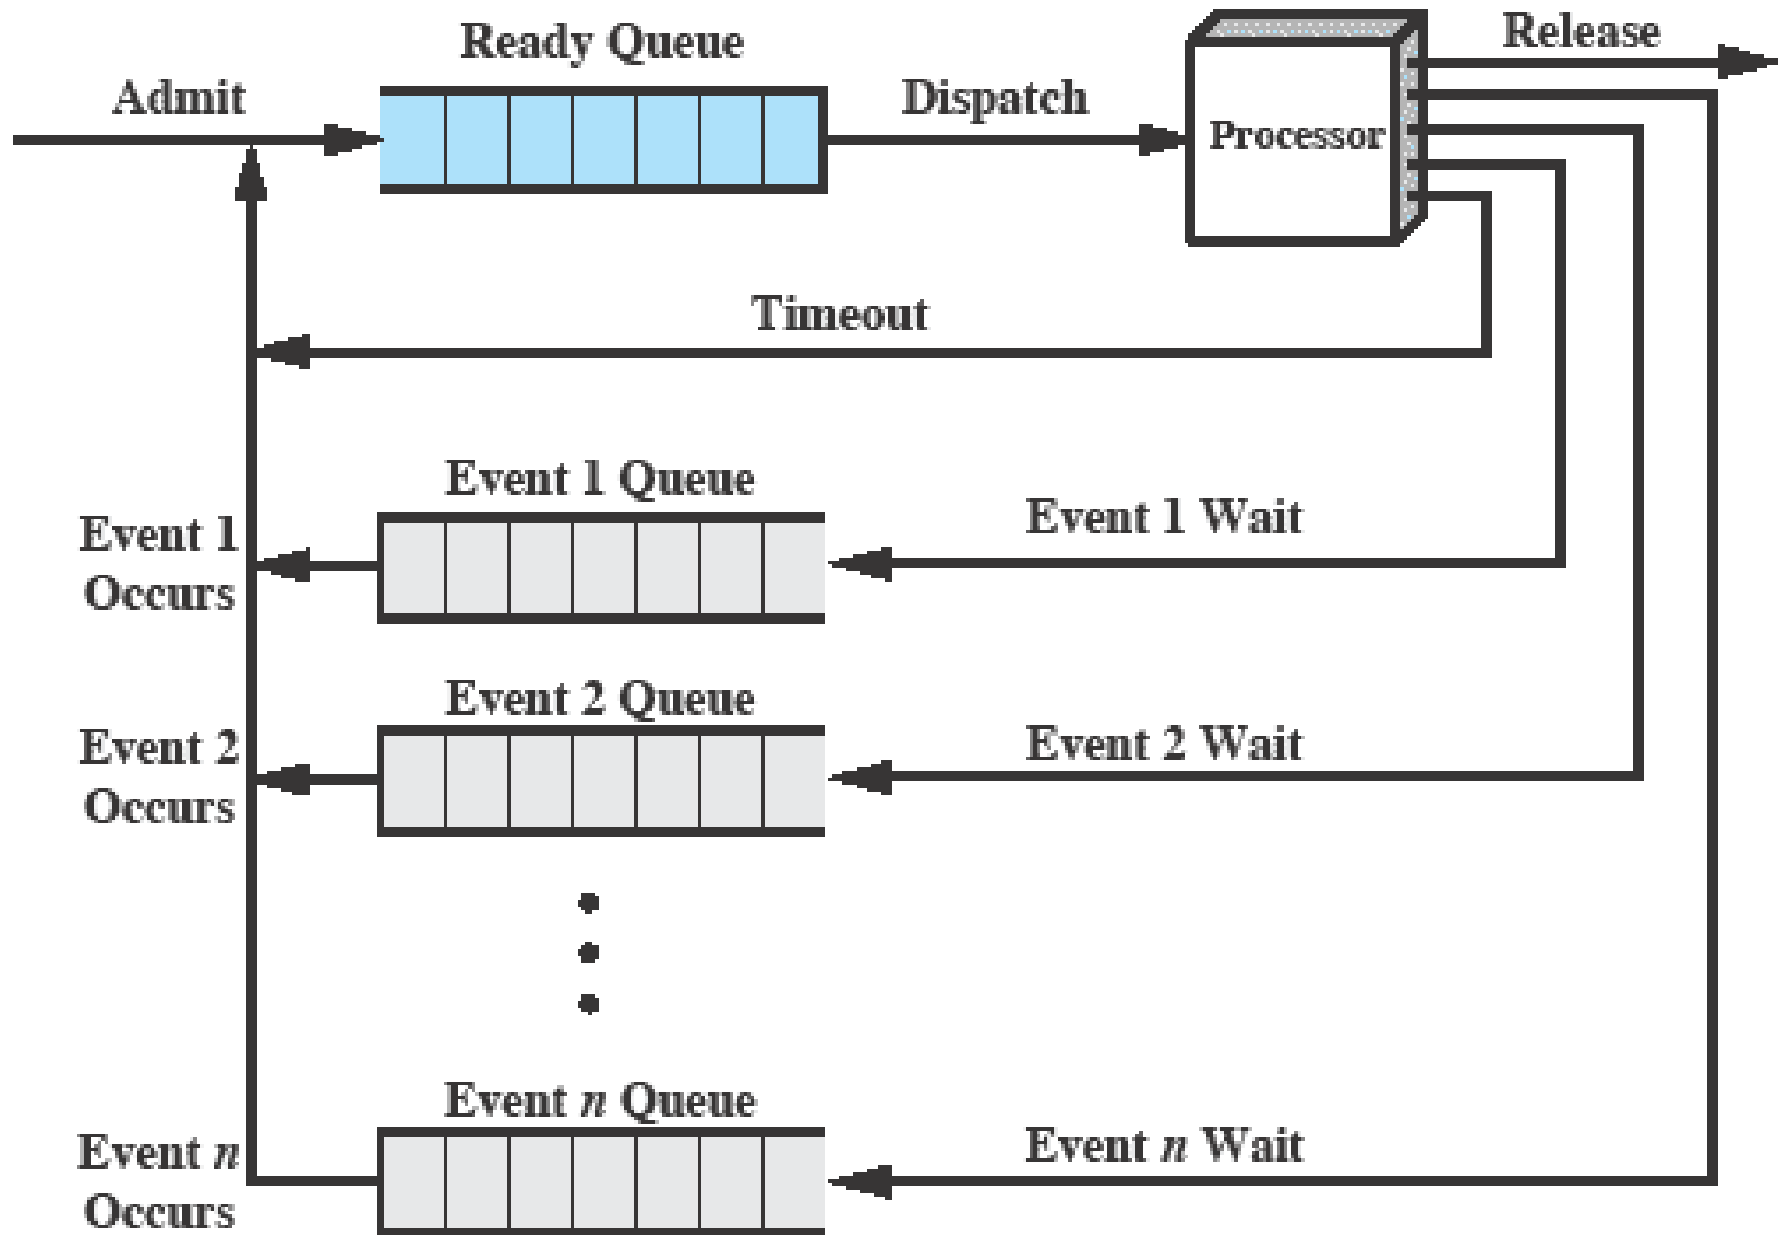
Process Creation

- .New batch job
- .Interactive logon
- .Created by OS to provide a service
- .Spawned by existing process – UNIX: fork + exec

Process Termination

- Normal completion
- Time limit exceeded
- Privileged instructions in user mode
- Parent termination
- Operator or OS intervention (e.g. deadlock)
- Errors
 - memory unavailable
 - bounds violation
 - protection error
 - arithmetic error
 - I/O failure

Multiple Blocked Queues



Suspended Processes

- Processor is faster than I/O so many processes could be waiting for I/O
- Swap these processes to disk to free up more memory
- Blocked state becomes suspend state when swapped to disk
- Two new states
 - Blocked/Suspend
 - Ready/Suspend

Two Suspend States



Operating System Control Structures

- Information about the current status of processes and resources
- Tables are constructed for each entity the operating system manages: memory table, I/O table, file table

Memory Tables

- Allocation of main memory to processes
- Allocation of secondary memory to processes
- Protection attributes for access to shared memory regions
- Information needed to manage virtual memory (page table)

I/O Tables

- I/O device is available or assigned
- Status of I/O operation
- Location in main memory being used as the source or destination of the I/O transfer

File Tables

- Existence of files
- Location on secondary memory
- Current status
- Attributes
- Sometimes this information is maintained by a file management system

OS Control Tables

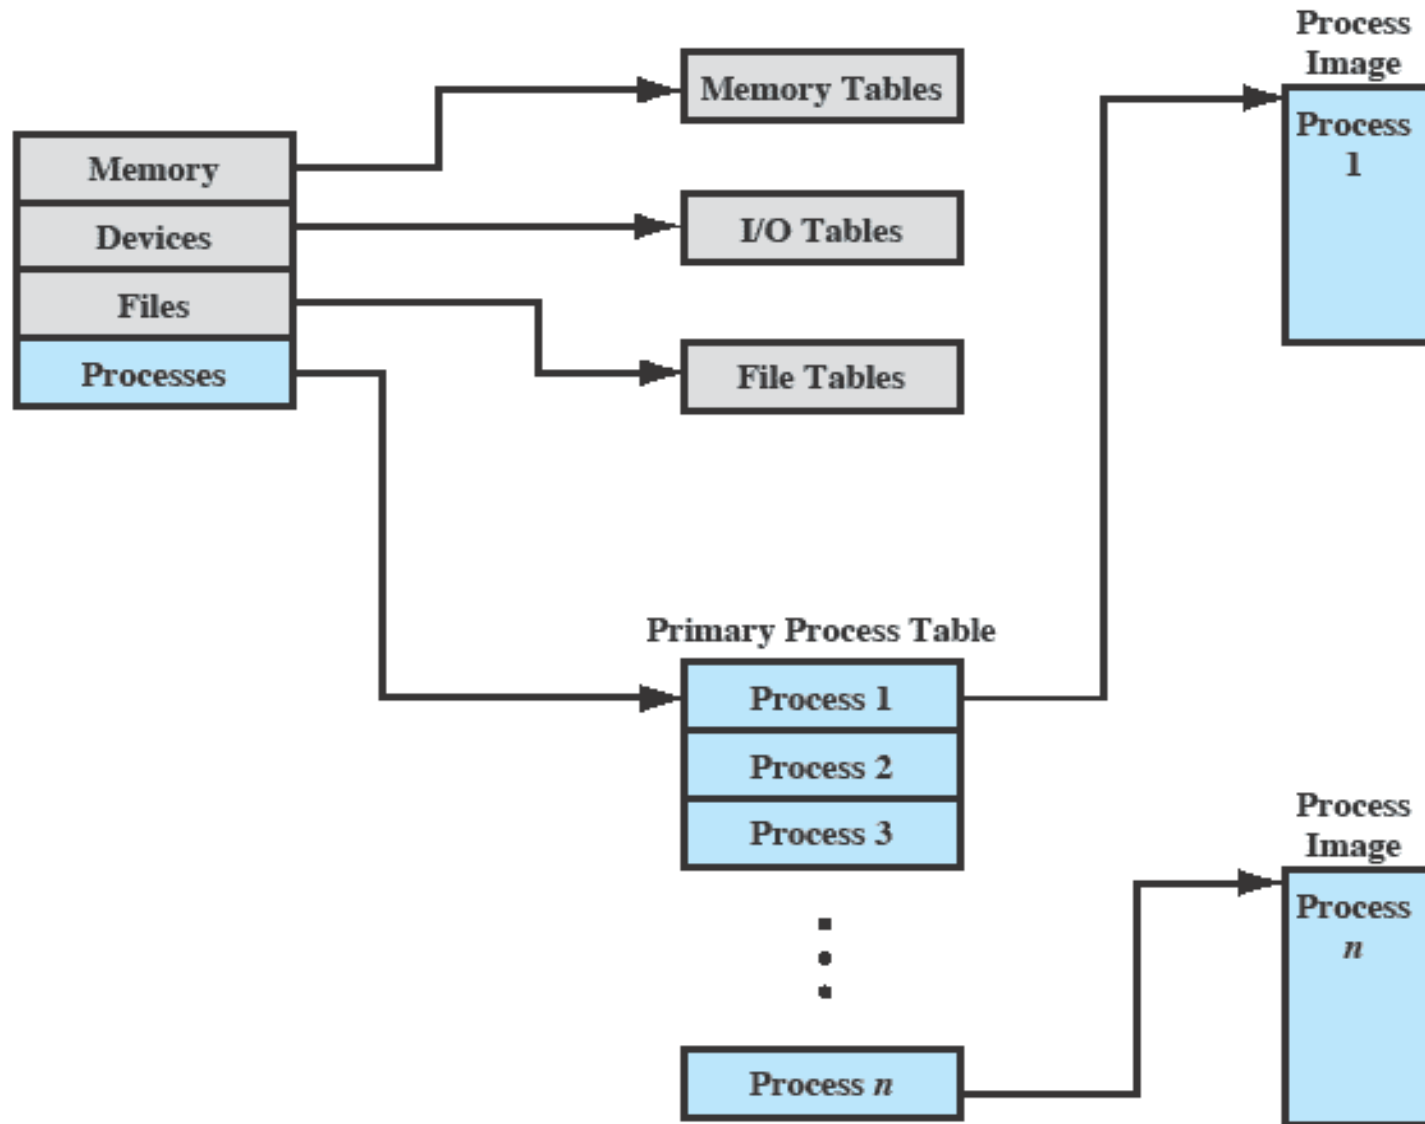


Figure 3.11 General Structure of Operating System Control Tables

Process Table – Process Control Block

- Identifiers: process id, parent process, user id
- User-visible registers
- Control and status registers: PC, PSW
- Stack pointer
- Scheduling and state info: process state, priority, used CPU time, event (waiting for)
- Process privileges
- Memory management: pointers to segments, page table
- Resource ownership and utilization

Modes of Execution

- User mode
 - Less-privileged mode
 - User programs typically execute in this mode
- System mode, control mode, or kernel mode
 - More-privileged mode
 - Kernel of the operating system
- Mode switch

Process Creation

- Assign a unique process identifier
- Allocate memory space for the process
- Initialize process control block
- Set up appropriate linkages (e.g. put process in scheduling queue)
- Create or expand other data structures (e.g. CPU time, page table)

When to Switch Process

Clock interrupt: process has executed for the maximum allowable time slice

I/O interrupt

Memory fault: memory address is in virtual memory so it must be brought into main memory – requires I/O

Trap: error or exception occurred; may cause process to be moved to Exit state

Supervisor/system call, e.g., such as file open

Change of Process State

1. Save context of processor including program counter and other registers
2. Update the process control blocks
3. Move process into appropriate queue – ready; blocked; ready/suspend
4. Run the scheduler to select another process for execution
5. Update the process control block
6. Restore context of the selected process

Execution of the Operating System

- Non-process Kernel

- Execute kernel outside of any process
- Operating system code is executed as a separate entity that operates in privileged mode – monolithic OS


- Execution within user processes

- Operating system software within context of a user process, e.g. scheduler


- Process-based operating system

- Implement the OS as a collection of system processes – modular OS

Execution of the Operating System



(a) Separate kernel



(b) OS functions execute within user processes



(c) OS functions execute as separate processes

Processes and Threads

- Resource ownership - process includes a virtual address space to hold the process image
- Scheduling/execution- follows an execution path that may be interleaved with other processes
- These two characteristics are treated independently by the operating system

Threads

Process =

resource grouping (code, data, open files,
etc.) +

execution (program counter, registers,
stack)

Multithreading:

- multiple execution takes place in the same process environment
- co-operation by sharing resources (address space, open files, etc.)

The Thread Model


Per process items	Per thread items
Address space Global variables Open files Child processes Pending alarms Signals and signal handlers Accounting information	Program counter Registers Stack State

Left: Items shared by all threads in a process

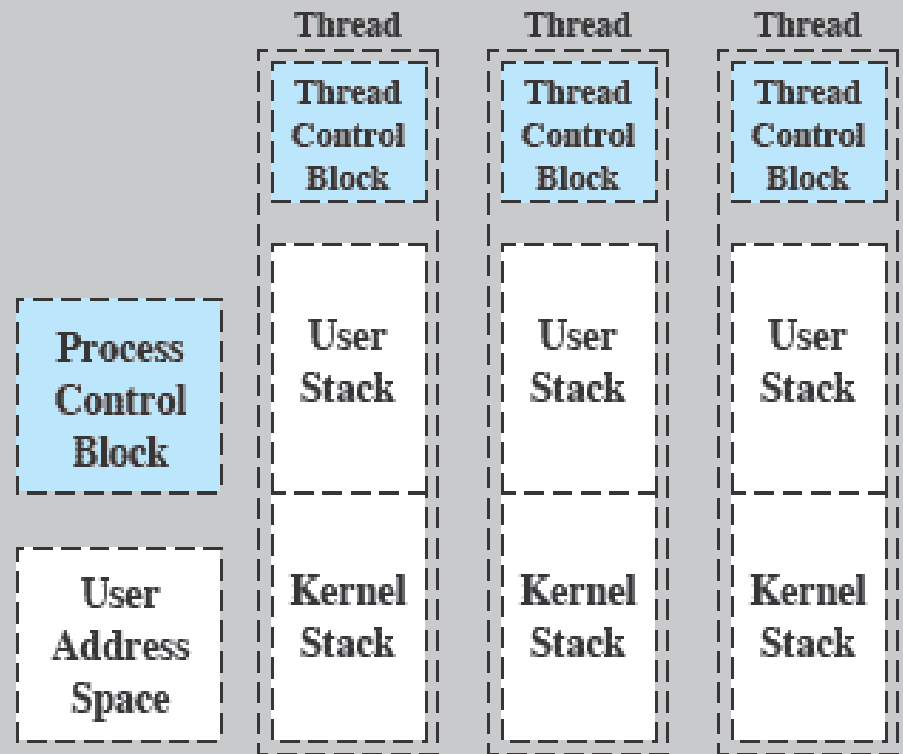
Right: Items private to each thread

Threads

Single-Threaded Process Model



Multithreaded Process Model



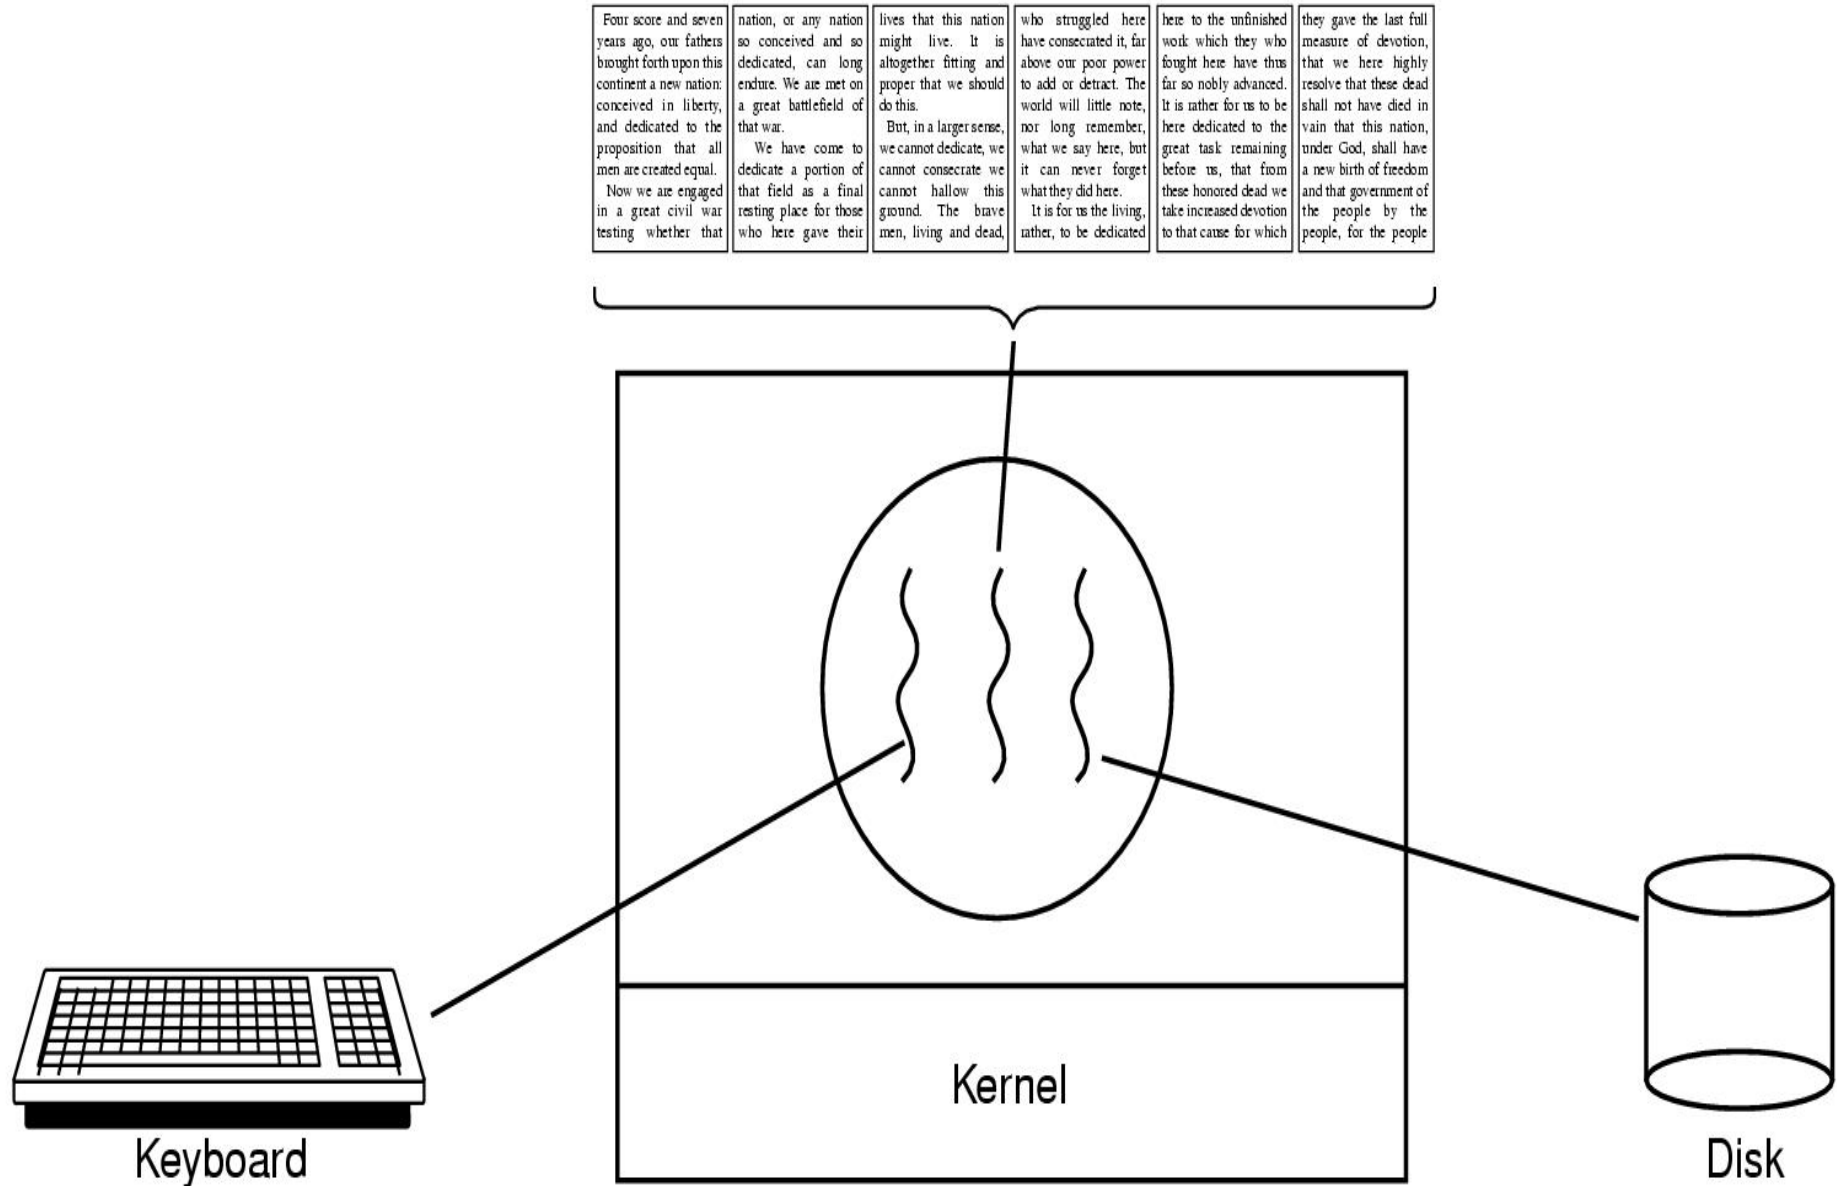
Benefits of Threads

- Takes less time to create a new thread than a process
- Less time to terminate a thread than a process
- Less time to switch between two threads within the same process
- Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel

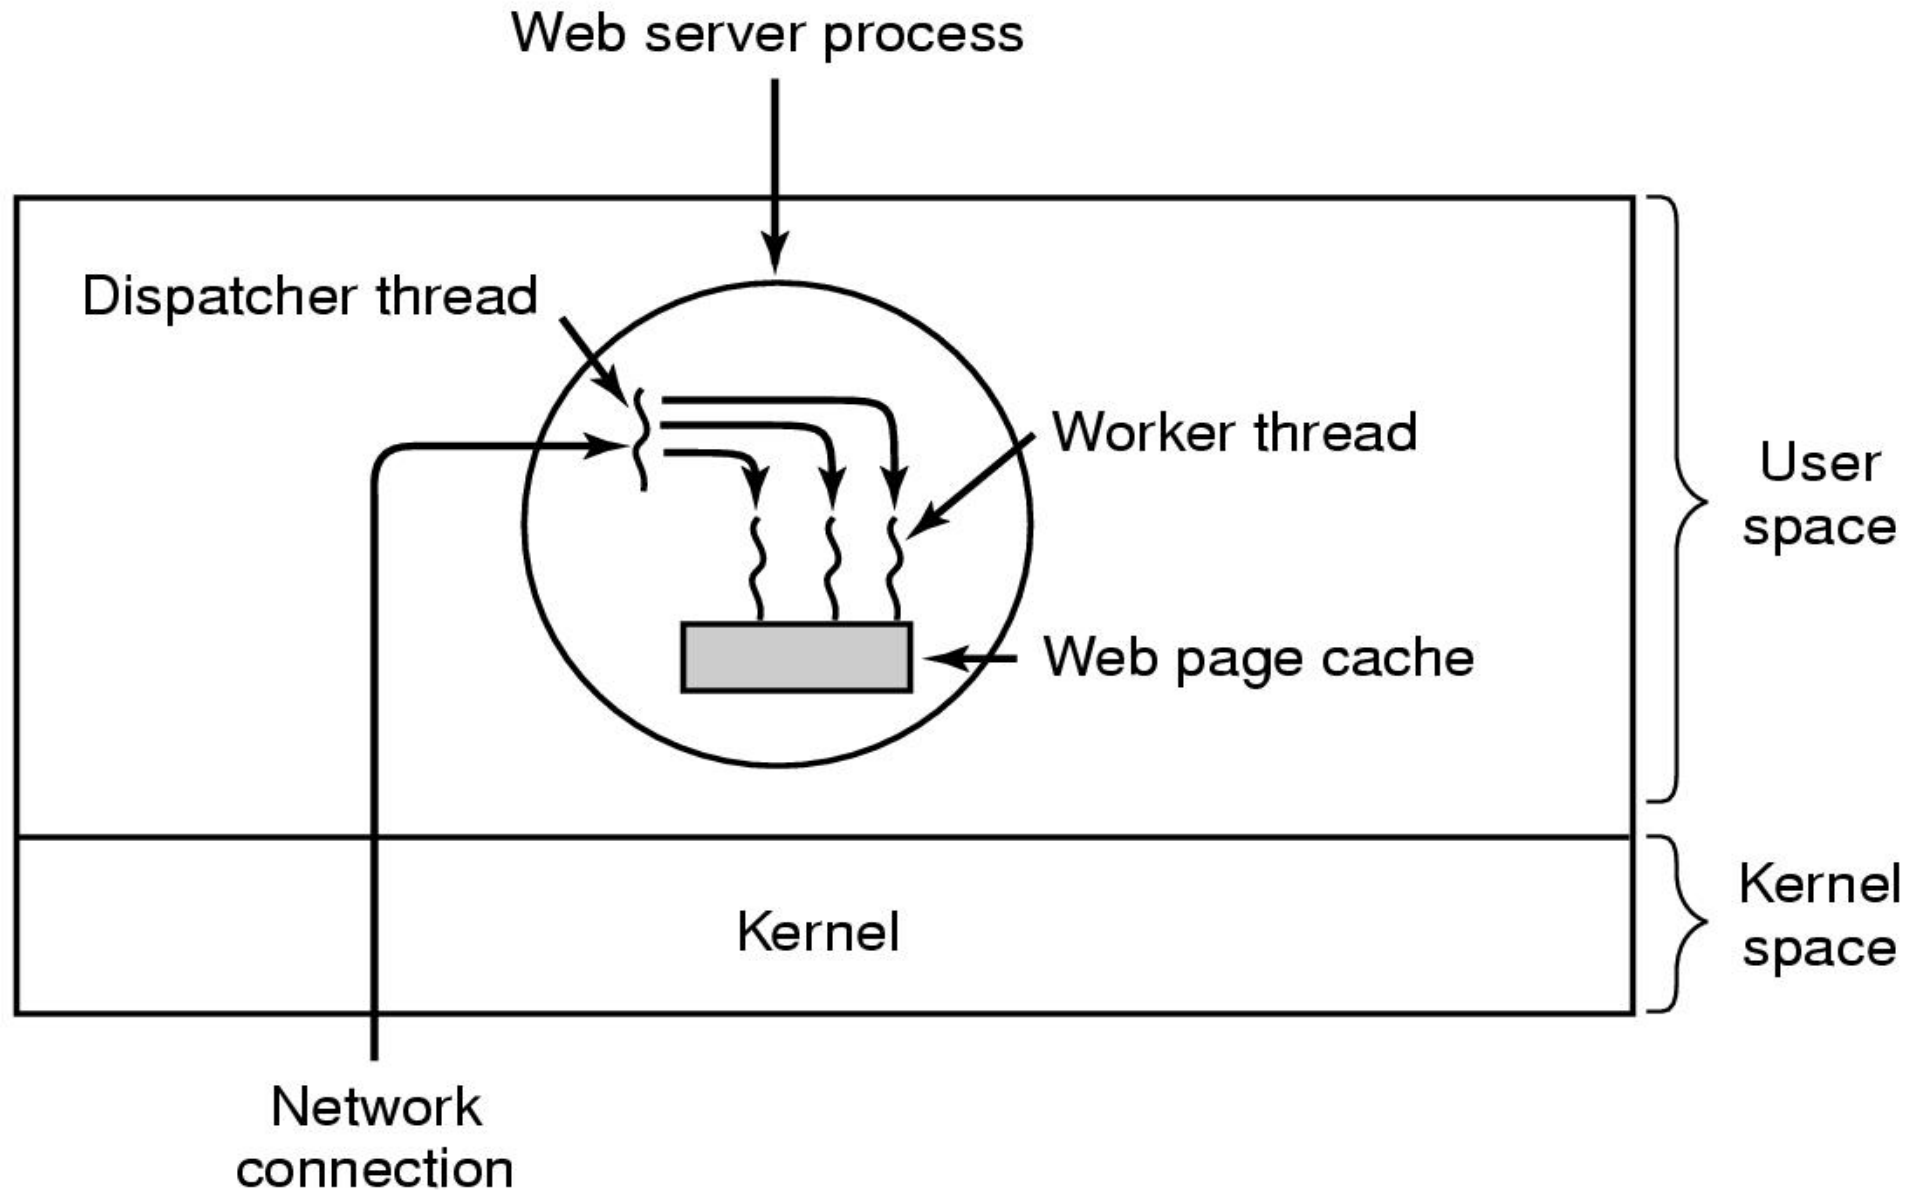
Uses of Threads in a Single-User Multiprocessing System

- Foreground and background work
- Speed of execution, e.g. blocked and running threads in one process
- Modular program structure
- Specific scheduling algorithms


A word processor with three threads



A multithreaded Web server



User-Level Threads



(a) Pure user-level


User-Level Threads

- All thread management is done by the application
- The kernel is not aware of the existence of threads
- Blocking system call!!!

Kernel-Level Threads

- Windows is an example of this approach
- Kernel maintains context information for the process and the threads
- Scheduling is done on a thread basis

Kernel-Level Threads




(b) **Pure kernel-level**

Combined Approaches

- Example is Solaris
- Thread creation done in the user space
- Bulk of scheduling and synchronization of threads within application

Combined Approach



Solaris

- Process includes the user's address space, stack, and process control block
- User-level threads
- Lightweight processes (LWP)
- Kernel threads

Processes and Threads in Solaris

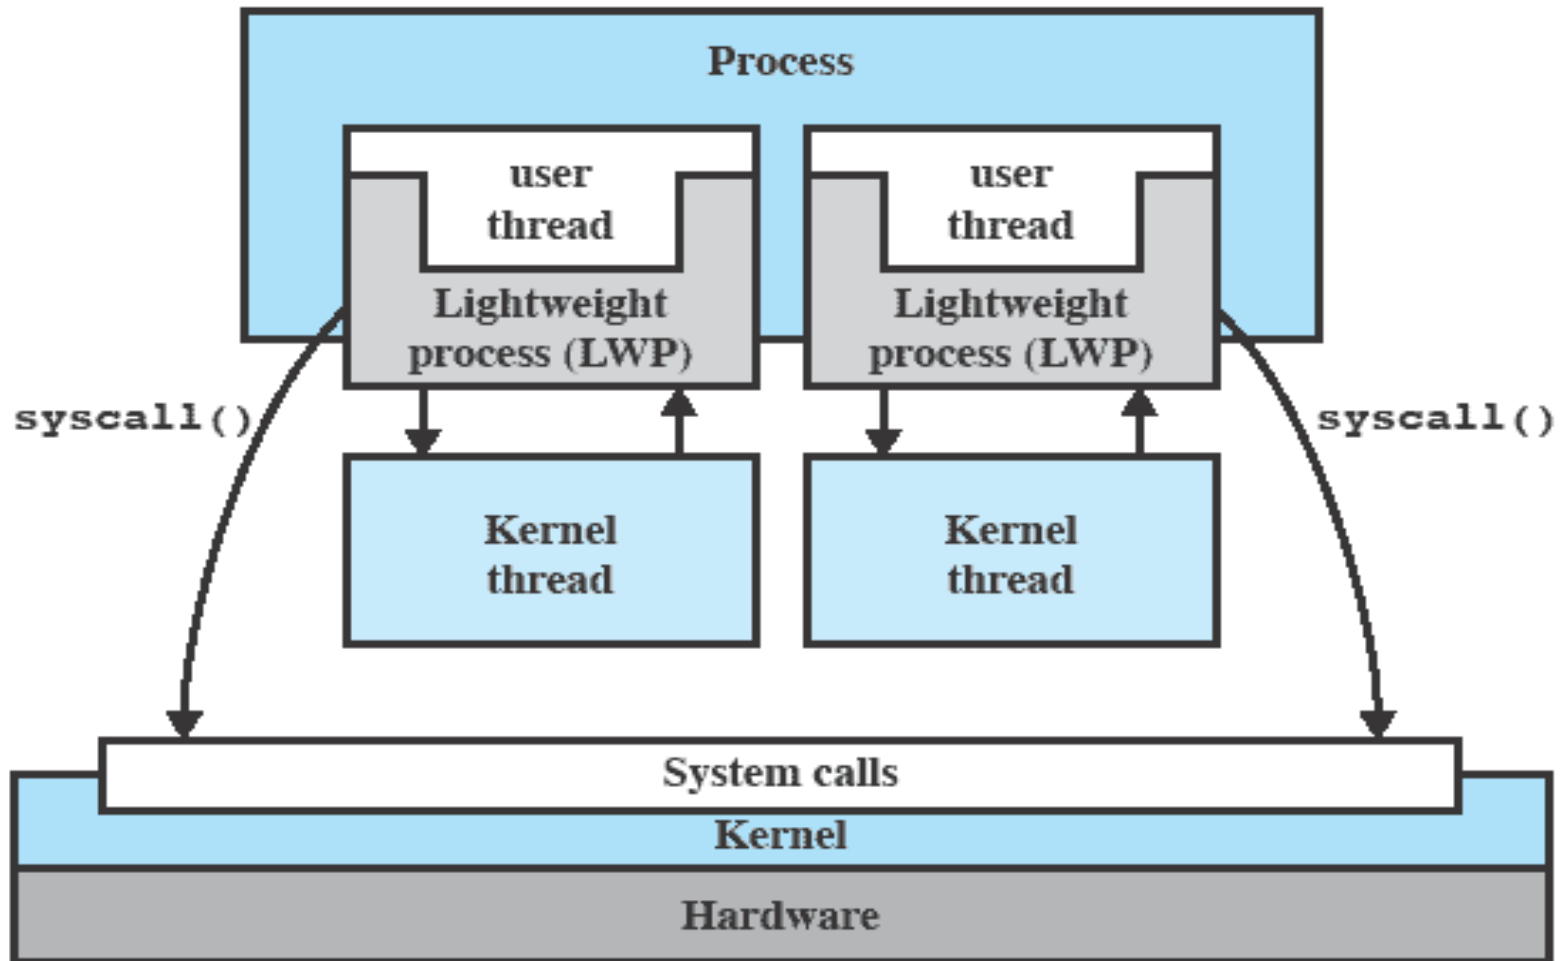


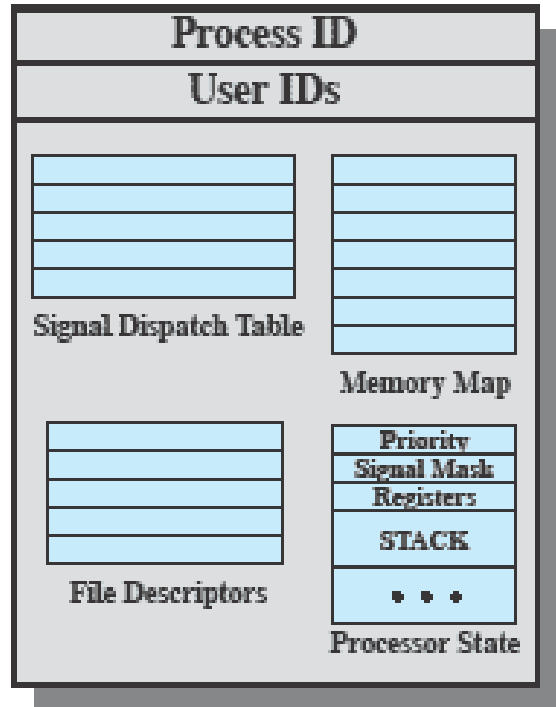
Figure 4.15 Processes and Threads in Solaris [MCDO07]

LWP Data Structure

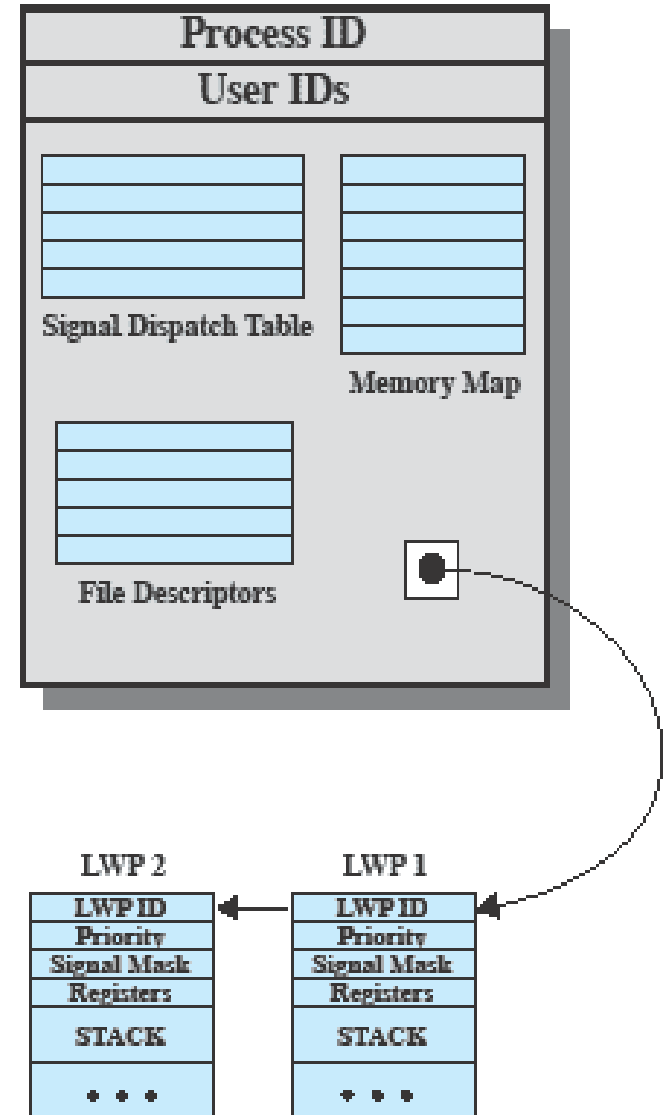
- Identifier
- Priority
- Signal mask
- Saved values of user-level registers
- Kernel stack
- Resource usage and profiling data
- Pointer to the corresponding kernel thread
- Pointer to the process structure

Process Structure

UNIX Process Structure




Solaris Process Structure



Windows Processes

- Implemented as objects
- An executable process may contain one or more threads
- Both processes and thread objects have built-in synchronization capabilities

Windows Processes



Windows Process Object - Job

Object Type

Process

**Object Body
Attributes**

Process ID
Security Descriptor
Base priority
Default processor affinity
Quota limits
Execution time
I/O counters
VM operation counters
Exception/debugging ports
Exit status

Services

Create process
Open process
Query process information
Set process information
Current process
Terminate process

Windows Thread Object

Object Type

Thread

**Object Body
Attributes**

Thread ID
Thread context
Dynamic priority
Base priority
Thread processor affinity
Thread execution time
Alert status
Suspension count
Impersonation token
Termination port
Thread exit status

Services

Create thread
Open thread
Query thread information
Set thread information
Current thread
Terminate thread
Get context
Set context
Suspend
Resume
Alert thread
Test thread alert
Register termination port

Thread States

