## Operating Systems: Internals and Design Principles, 6/E William Stallings

Chapter 11 I/O Management and Disk Scheduling

Patricia Roy
Manatee Community College, Venice,
FL
©2008, Prentice Hall

## Differences in I/O Devices

#### Data rate

- May be differences of several orders of magnitude between the data transfer rates
- Unit of transfer Character and Block devices
  - Data may be transferred as a stream of bytes for a terminal or in larger blocks for a disk
- Data representation
  - Encoding and error-correction schemes
- Error conditions
  - Different types of errors

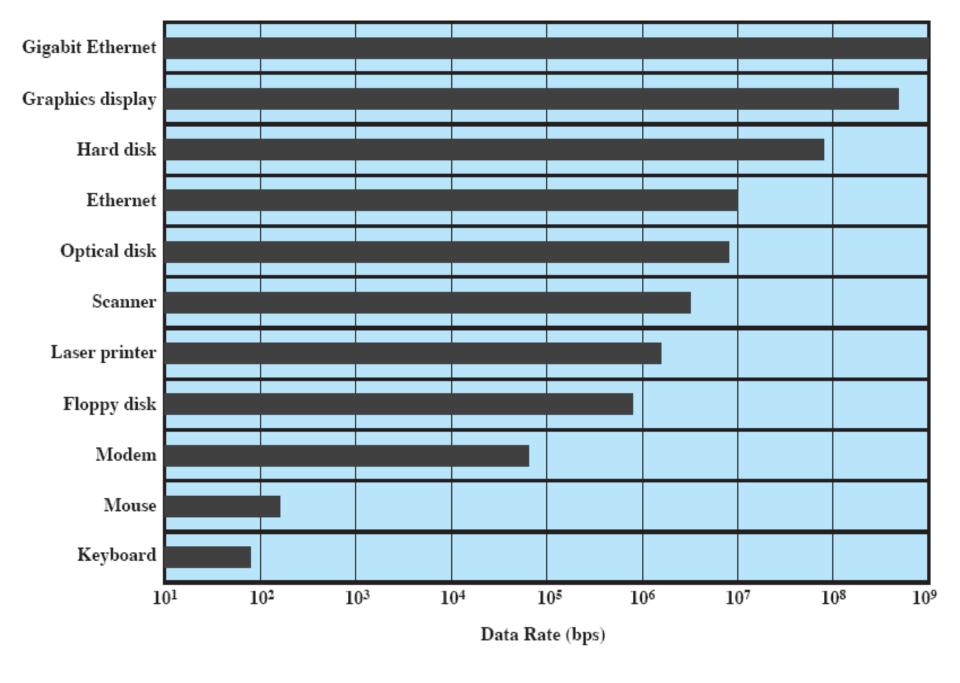


Figure 11.1 Typical I/O Device Data Rates

### **Device Controllers**

- I/O devices have components:
  - mechanical component
  - electronic component
- The electronic component is the device controller
  - may be able to handle multiple devices
- Controller's tasks
  - convert serial bit stream to block of bytes
  - perform error correction as necessary
  - communicate with CPU

# Performing I/O

Programmed I/O

Process(or) is busy-waiting for the operation to complete

• Interrupt-driven I/O

I/O command is issued

Processor continues executing other instructions

Direct Memory Access

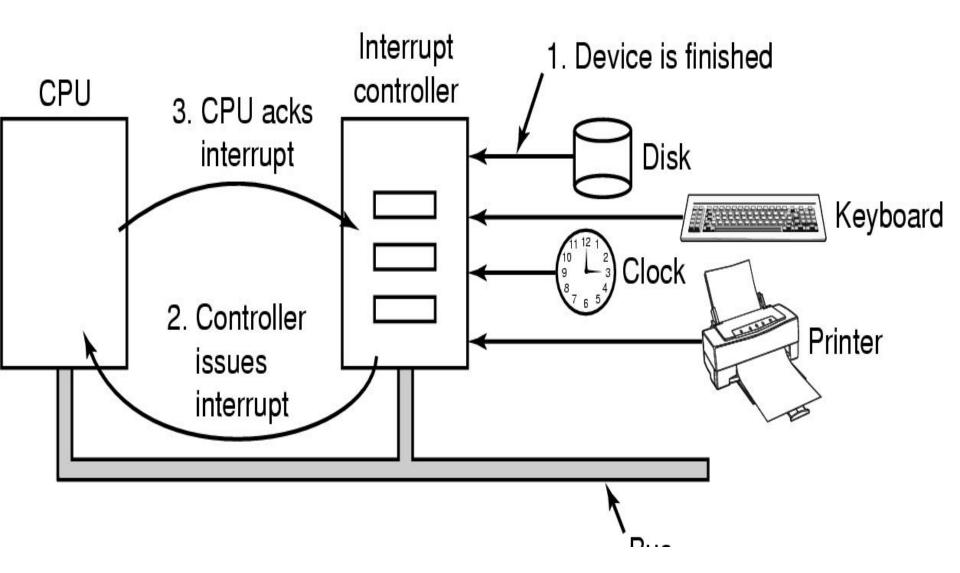
DMA module controls exchange of data between main memory and the I/O device

Processor interrupted only after entire I/O finished

# Programmed I/O

Writing a string to the printer using programmed I/O --- busy waiting

# Interrupts Revisited



# Interrupt-Driven I/O

```
copy from user(buffer, p, count);
                                             if (count == 0) {
                                                 unblock user();
enable interrupts();
while (*printer_status_reg != READY);
                                             } else {
*printer_data_register = p[0];
                                                 *printer_data_register = p[i];
scheduler();
                                                 count = count - 1;
                                                  i = i + 1;
                                             acknowledge interrupt();
                                             return_from_interrupt();
```

• Writing a string to the printer using interrupt-driven I/O

(b)

- Code executed when print system call is made
- Interrupt service procedure

(a)

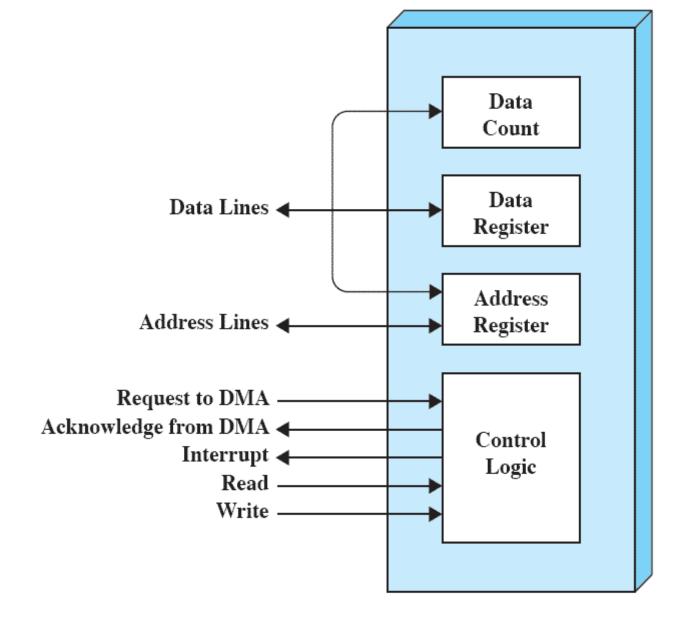
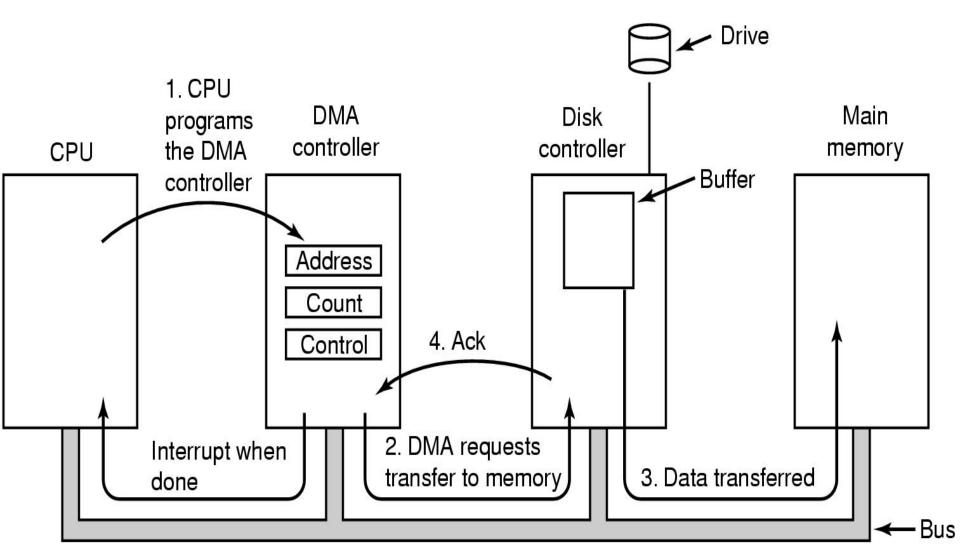
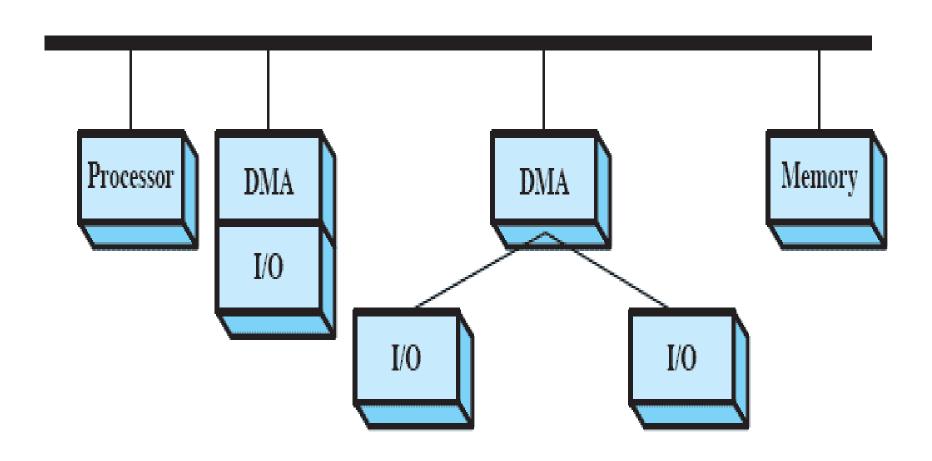


Figure 11.2 Typical DMA Block Diagram

# Direct Memory Access (DMA)

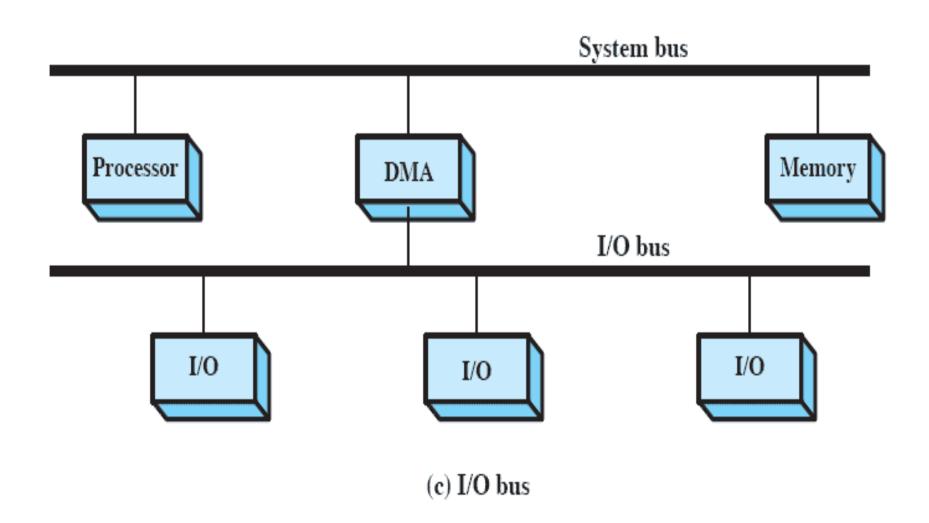


# DMA Configurations (1)



(b) Single-bus, Integrated DMA-I/O

# DMA Configurations (2)



# I/O Using DMA

(a) (b)

- Printing a string using DMA
  - code executed when the print system call is made
  - interrupt service procedure

#### Table 11.1 I/O Techniques

	No Interrupts	Use of Interrupts
I/O-to-memory transfer through processor	Programmed I/O	Interrupt-driven I/O
Direct I/O-to-memory transfer		Direct memory access (DMA)

## Operating System Design Issues (1)

#### Efficiency

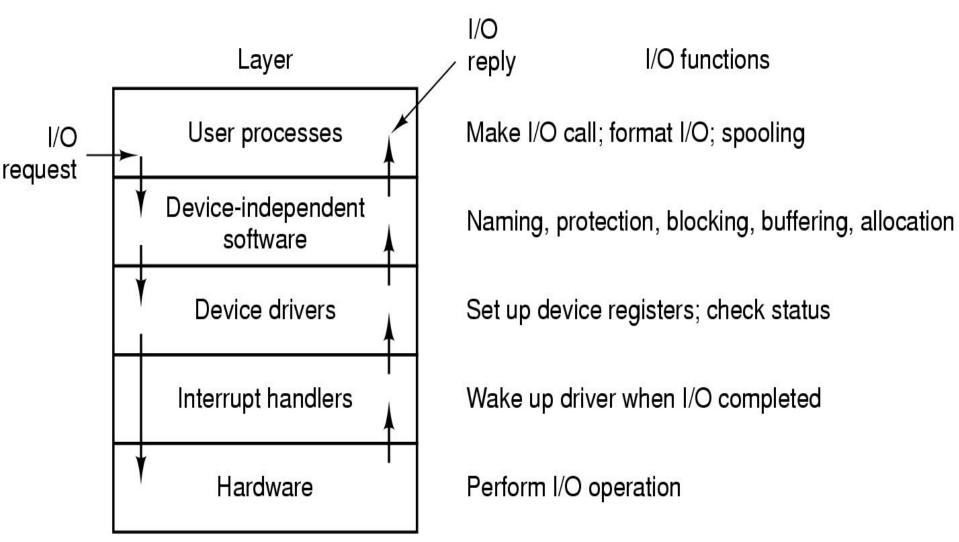
- Most I/O devices extremely slow compared to main memory
- Use of multiprogramming allows for some processes to be waiting on I/O while another process executes
- I/O cannot keep up with processor speed
- Swapping is used to bring in additional Ready processes which is an I/O operation

## Operating System Design Issues (2)

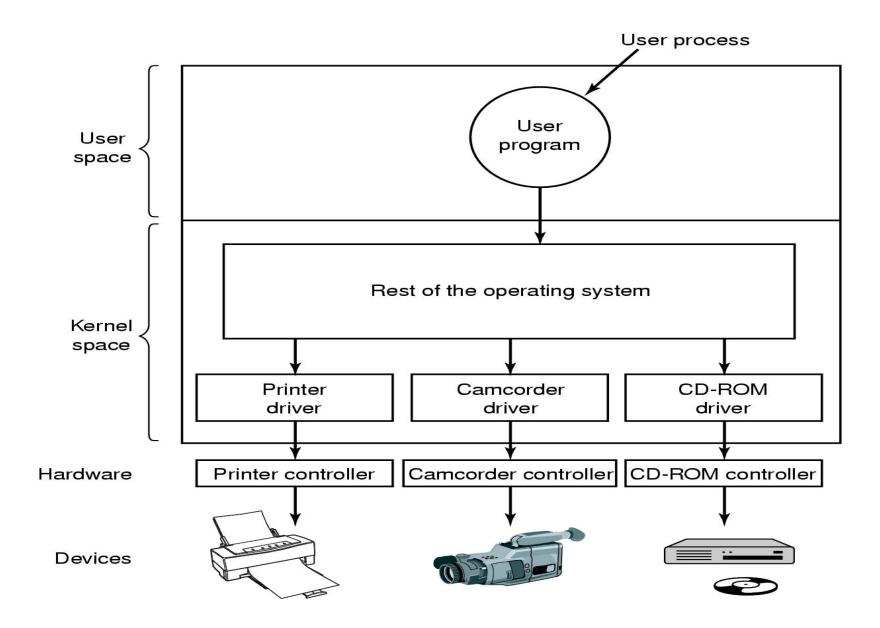
#### Generality

- Desirable to handle all I/O devices in a uniform manner
- Hide most of the details of device I/O in lowerlevel routines

# I/O Software Layers



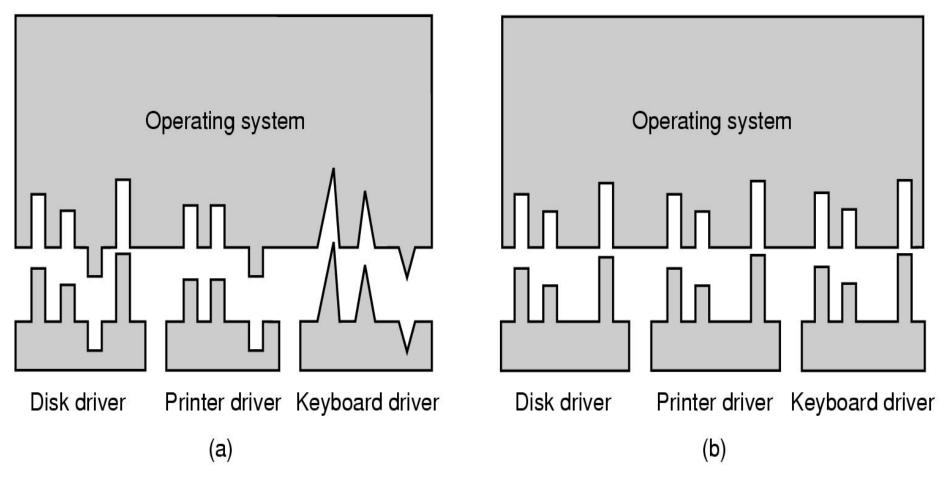
## **Device Drivers**



### Tasks of Device Drivers

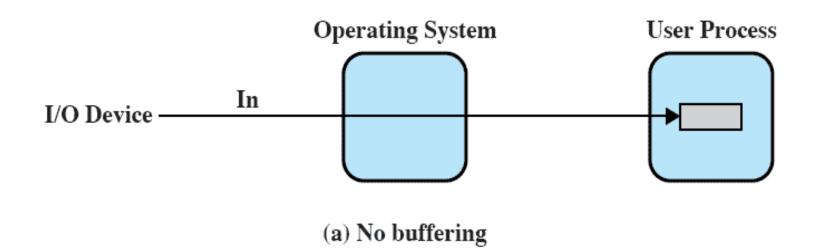
- Accept abstract requests
- Check input parameters
- Translate from abstract to concrete
- Check if device is in use
- Issue commands to controller
- (Block)
- Check errors
- Return (error) to caller

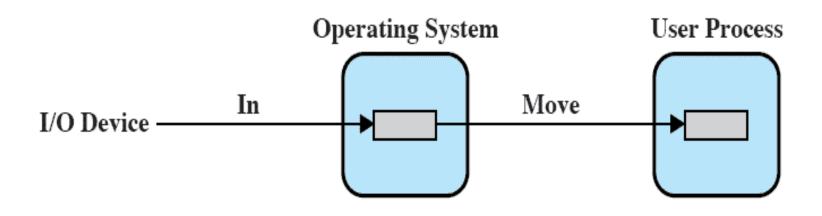
# Device-Independent I/O Software



- (a) Without a standard driver interface
- (b) With a standard driver interface

# With or without Buffering





(b) Single buffering

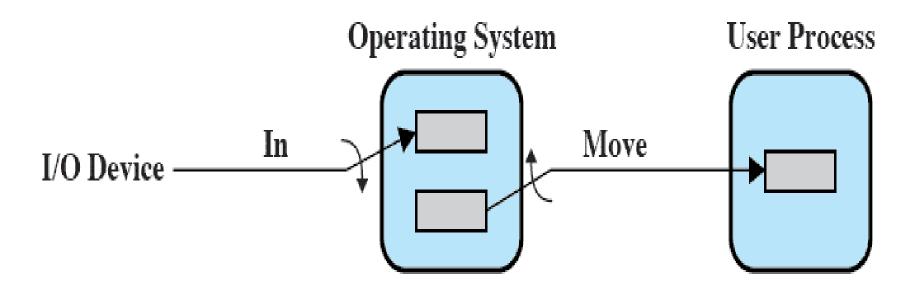
## Single Buffer

- Operating system assigns a buffer in main memory for an I/O request
  - Input transfers made to buffer
  - Data moved to user space when needed
  - Extra data is moved into the buffer
  - User process can process data while additional data is read in
  - Swapping can occur since input is taking place in system memory, not user memory
  - Operating system keeps track of assignment of system buffers to user processes

### Double Buffer

Use two system buffers instead of one

A process can transfer data to or from one buffer while the operating system empties or fills the other buffer

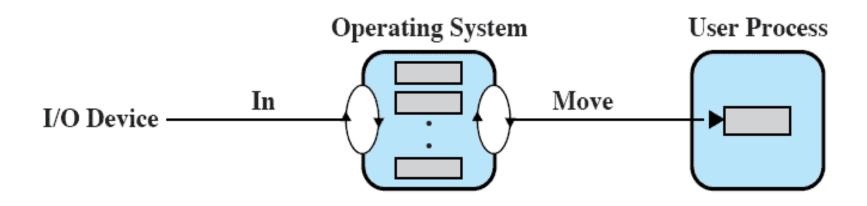


(c) Double buffering

## Circular Buffer

- More than two buffers are used
- Each individual buffer is one unit in a circular buffer

Used when I/O operation must keep up with process



(d) Circular buffering

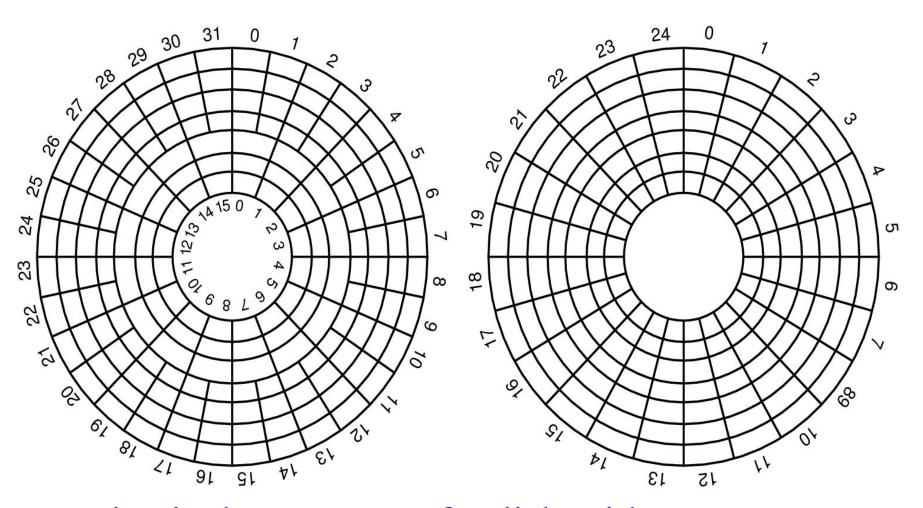
## Disk Performance Parameters (1)

- To read or write, the disk head must be positioned at the desired track and at the beginning of the desired sector
- Seek time
  - Time it takes to position the head at the desired track
- Rotational delay or rotational latency
  - Time it takes for the beginning of the sector to reach the head

## Disk Performance Parameters (2)

- Access time
  - Sum of seek time and rotational delay
  - The time it takes to get in position to read or write
- Data transfer occurs as the sector moves under the head

### Disk Hardware



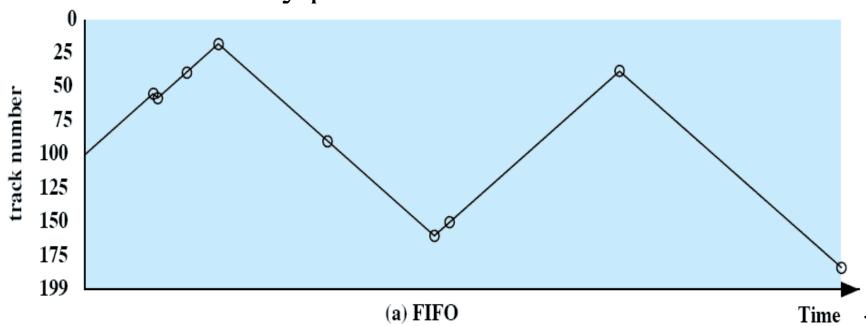
- Physical geometry of a disk with two zones
- A possible virtual geometry for this disk

# Disk Scheduling Policies

- Seek time is the reason for differences in performance
- For a single disk there will be a number of I/O requests
- If requests are selected randomly, get poor performance

# Disk Scheduling Policies - FIFO

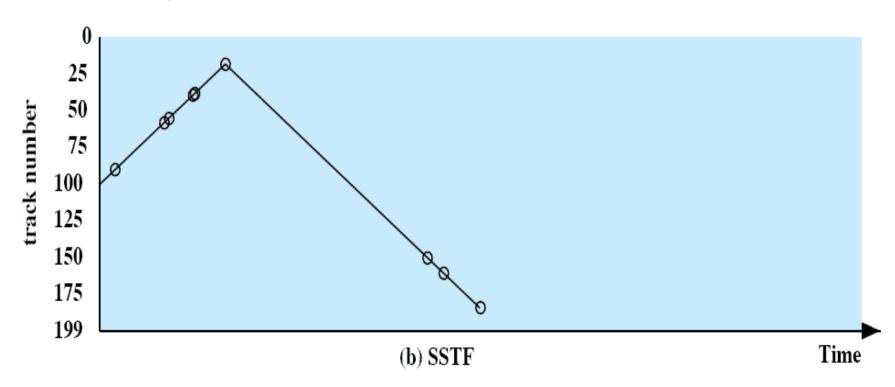
- First-in, first-out (FIFO)
  - Process request sequentially
  - Fair to all processes
  - •Approaches random scheduling in performance if there are many processes



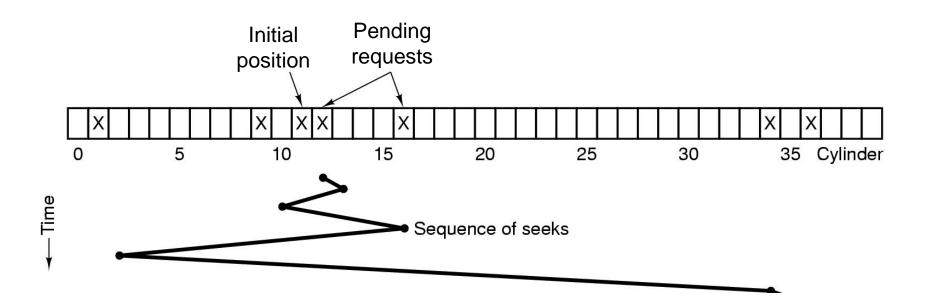
29

# Disk Scheduling Policies - SSTF

- Shortest Service/Seek Time First
  - Select the disk I/O request that requires the least movement of the disk arm from its current position
  - •Always choose the minimum seek time



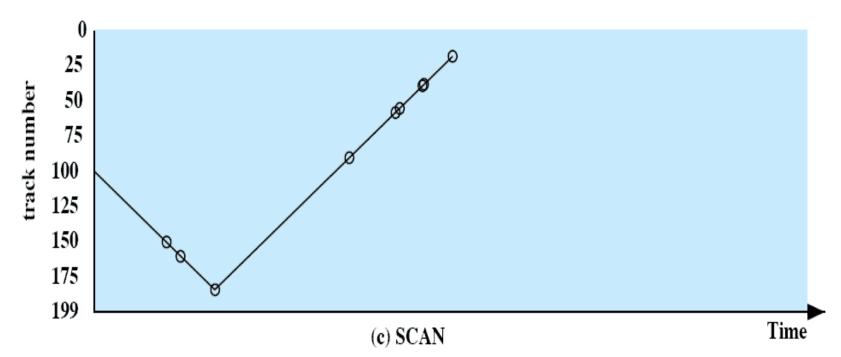
# SSF (2)



Shortest Seek First (SSF) disk scheduling algorithm

# Disk Scheduling Policies - SCAN

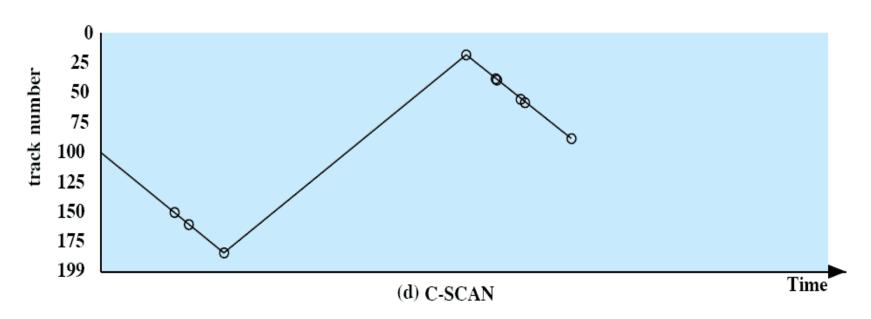
- SCAN or Elevator
  - Arm moves in one direction only, satisfying all outstanding requests until it reaches the last track in that direction
  - Direction is reversed



## Disk Scheduling Policies – C-SCAN

#### C-SCAN

- Restricts scanning to one direction only
- •When the last track has been visited in one direction, the arm is returned to the opposite end of the disk and the scan begins again

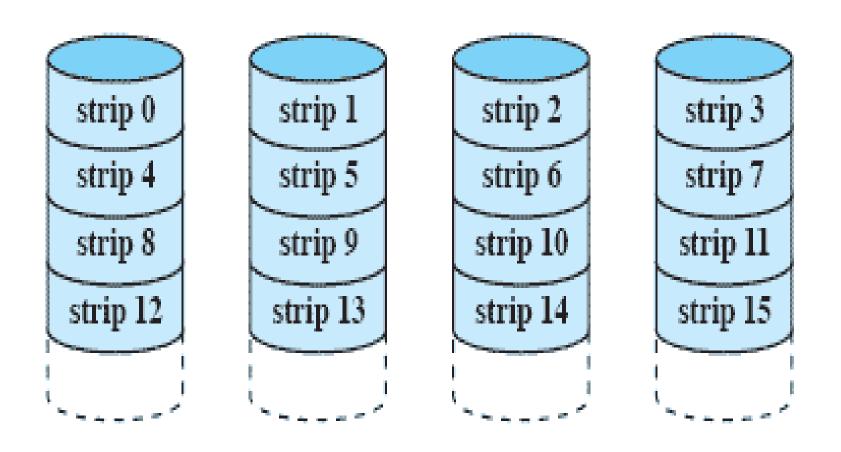


(a)	FIFO	(b)	SSTF	(c) S	SCAN	(d) C-	SCAN
(starting at track 100)		(starting at track 100)		(starting at track 100, in the direction of increasing track number)		(starting at track 100, in the direction of increasing track number)	
Next track accessed	Number of tracks traversed	Next track accessed	Number of tracks traversed	Next track accessed	Number of tracks traversed	Next track accessed	Number of tracks traversed
55	45	90	10	150	50	150	50
58	3	58	32	160	10	160	10
39	19	55	3	184	24	184	24
18	21	39	16	90	94	18	166
90	72	38	1	58	32	38	20
160	70	18	20	55	3	39	1
150	10	150	132	39	16	55	16
38	112	160	10	38	1	58	3
184	146	184	24	18	20	90	32
Average seek length	55.3	Average seek length	27.5	Average seek length	27.8	Average seek length	35.8

#### **RAID**

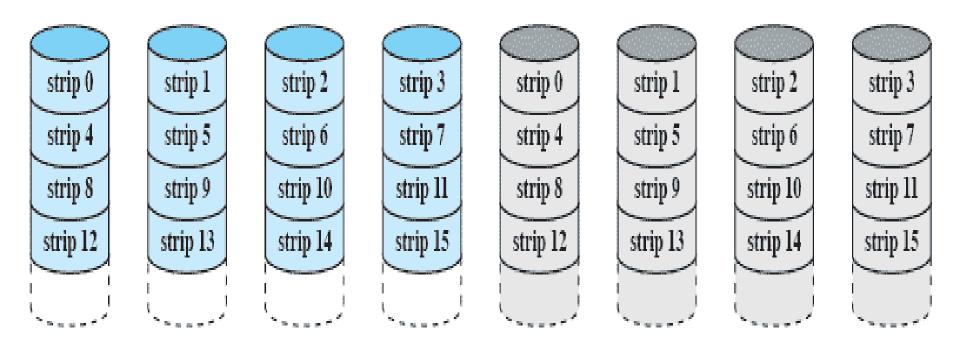
- Redundant Array of Independent Disks
- Set of physical disk drives viewed by the operating system as a single logical drive
- Data are distributed across the physical drives of an array
- Redundant disk capacity is used to store parity information

## RAID 0 (non-redundant)



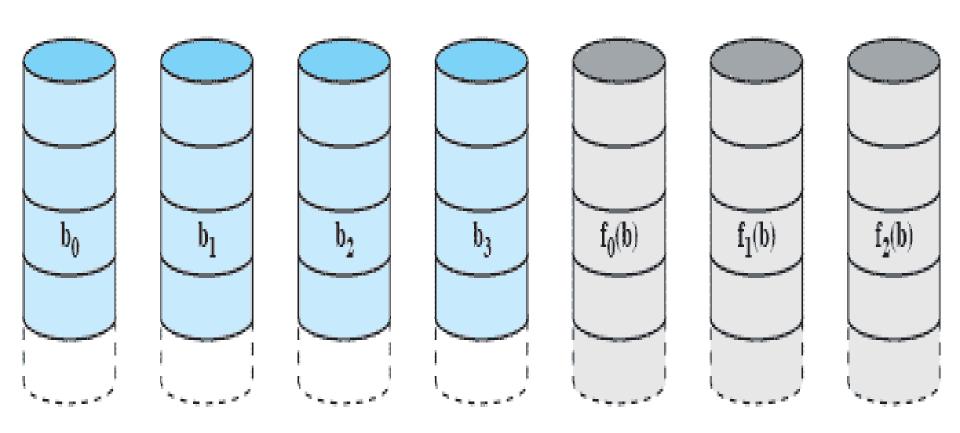
(a) RAID 0 (non-redundant)

# RAID 1 (mirrored)



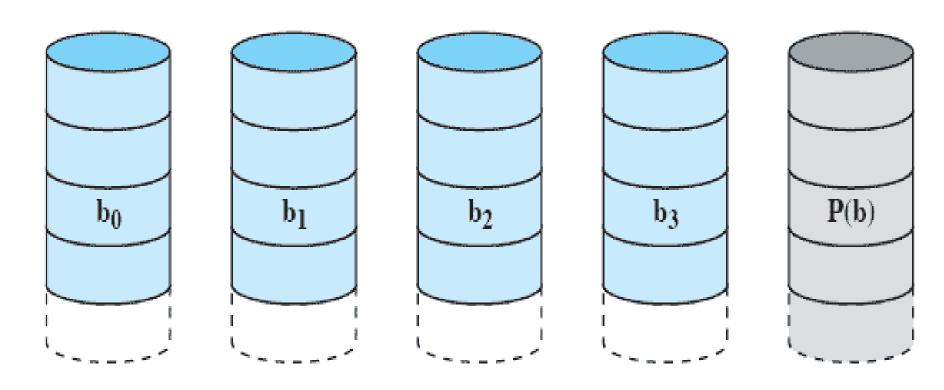
(b) RAID 1 (mirrored)

# RAID 2 (Hamming code)



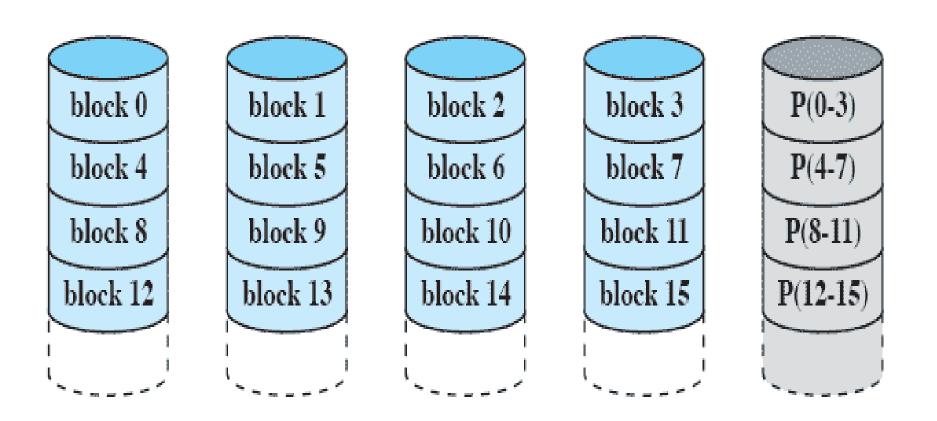
(c) RAID 2 (redundancy through Hamming code)

# RAID 3 (bit-interleaved parity)



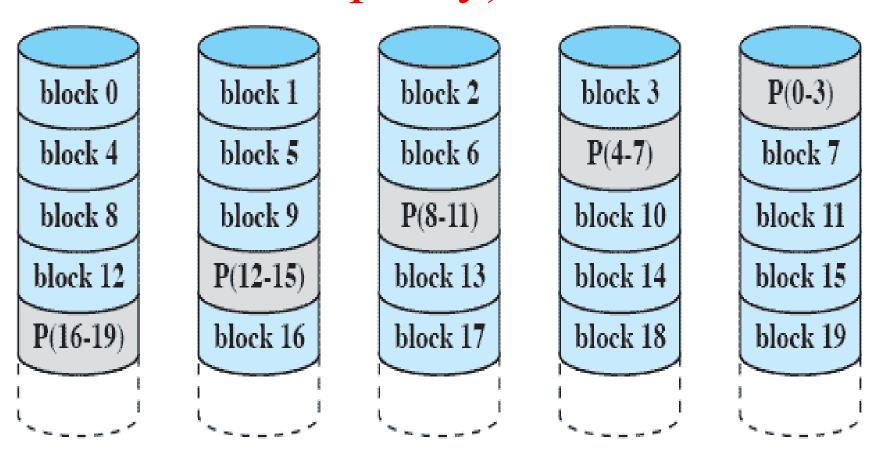
(d) RAID 3 (bit-interleaved parity)

# RAID 4 (block-level parity)



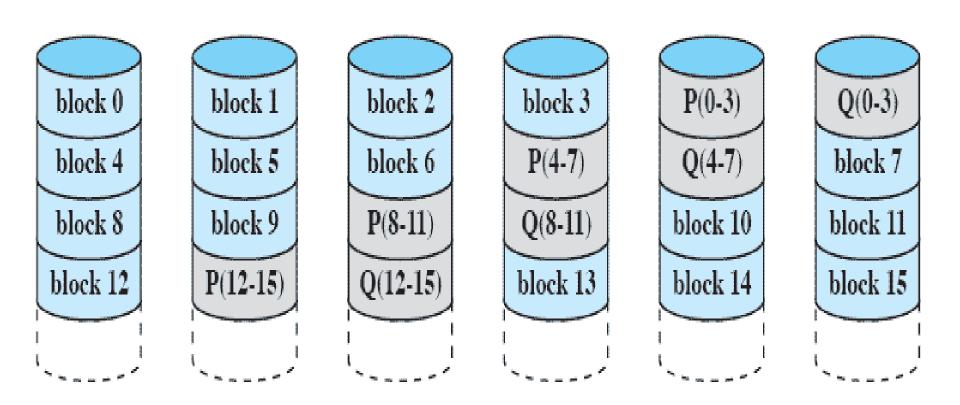
(e) RAID 4 (block-level parity)

# RAID 5 (block-level distributed parity)



(f) RAID 5 (block-level distributed parity)

## RAID 6 (dual redundancy)



(g) RAID 6 (dual redundancy)

#### Disk Cache

- Buffer in main memory for disk sectors
- Contains a copy of some of the sectors on the disk

# Least Recently Used (1)

- The block that has been in the cache the longest with no reference to it is replaced
- The cache consists of a stack of blocks
- Most recently referenced block is on the top of the stack
- When a block is referenced or brought into the cache, it is placed on the top of the stack

# Least Recently Used (2)

- The block on the bottom of the stack is removed when a new block is brought in
- Blocks don't actually move around in main memory
- A stack of pointers is used

# Least Frequently Used

- The block that has experienced the fewest references is replaced
- A counter is associated with each block
- Counter is incremented each time block accessed
- Block with smallest count is selected for replacement
- Some blocks may be referenced many times in a short period of time and the reference count is misleading

# Frequency-Based Replacement

